

```
In [1]: import numpy as np
import pandas as pd
import seaborn as sns
import matplotlib.pyplot as plt
from sklearn import preprocessing, svm
from sklearn.model_selection import train_test_split
from sklearn.linear_model import LinearRegression
from sklearn.preprocessing import StandardScaler
```

```
In [2]: df=pd.read_csv(r"C:\Users\sweet\Downloads\insuranceex.csv")
df
```

```
Out[2]:
```

	age	sex	bmi	children	smoker	region	charges
0	19	female	27.900	0	yes	southwest	16884.92400
1	18	male	33.770	1	no	southeast	1725.55230
2	28	male	33.000	3	no	southeast	4449.46200
3	33	male	22.705	0	no	northwest	21984.47061
4	32	male	28.880	0	no	northwest	3866.85520
...
1333	50	male	30.970	3	no	northwest	10600.54830
1334	18	female	31.920	0	no	northeast	2205.98080
1335	18	female	36.850	0	no	southeast	1629.83350
1336	21	female	25.800	0	no	southwest	2007.94500
1337	61	female	29.070	0	yes	northwest	29141.36030

1338 rows × 7 columns

Data Cleaning and Preprocessing

In [3]: `df.head()`

Out[3]:

	age	sex	bmi	children	smoker	region	charges
0	19	female	27.900	0	yes	southwest	16884.92400
1	18	male	33.770	1	no	southeast	1725.55230
2	28	male	33.000	3	no	southeast	4449.46200
3	33	male	22.705	0	no	northwest	21984.47061
4	32	male	28.880	0	no	northwest	3866.85520

In [4]: `df.tail()`

Out[4]:

	age	sex	bmi	children	smoker	region	charges
1333	50	male	30.97	3	no	northwest	10600.5483
1334	18	female	31.92	0	no	northeast	2205.9808
1335	18	female	36.85	0	no	southeast	1629.8335
1336	21	female	25.80	0	no	southwest	2007.9450
1337	61	female	29.07	0	yes	northwest	29141.3603

In [5]: `df.shape`

Out[5]: (1338, 7)

In [6]: `df.describe()`

Out[6]:


	age	bmi	children	charges
count	1338.000000	1338.000000	1338.000000	1338.000000
mean	39.207025	30.663397	1.094918	13270.422265
std	14.049960	6.098187	1.205493	12110.011237
min	18.000000	15.960000	0.000000	1121.873900
25%	27.000000	26.296250	0.000000	4740.287150
50%	39.000000	30.400000	1.000000	9382.033000
75%	51.000000	34.693750	2.000000	16639.912515
max	64.000000	53.130000	5.000000	63770.428010

In [7]: `df.info()`

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 1338 entries, 0 to 1337
Data columns (total 7 columns):
#   Column      Non-Null Count  Dtype
---  -
0   age         1338 non-null   int64
1   sex         1338 non-null   object
2   bmi         1338 non-null   float64
3   children    1338 non-null   int64
4   smoker      1338 non-null   object
5   region      1338 non-null   object
6   charges     1338 non-null   float64
dtypes: float64(2), int64(2), object(3)
memory usage: 73.3+ KB
```

```
In [8]: sns.lmplot(x="age",y="charges",data=df,order=2,ci=None)
```

```
Out[8]: <seaborn.axisgrid.FacetGrid at 0x12592b72250>
```



```
In [9]: df.fillna(method='ffill',inplace=True)
```


```
In [10]: x=np.array(df['age']).reshape(-1,1)
         y=np.array(df['charges']).reshape(-1,1)
```

```
In [11]: df.dropna(inplace=True)
```

```
In [12]: x_train,x_test,y_train,y_test=train_test_split(x,y,test_size=0.25)
         regr=LinearRegression()
         regr.fit(x_train,y_train)
         print(regr.score(x_test,y_test))
```


0.10968904282635228

```
In [13]: y_pred=regr.predict(x_test)
plt.scatter(x_test,y_test,color='g')
plt.plot(x_test,y_pred,color='y')
plt.show()
```




```
In [14]: df500=df[:][:500]  
sns.lmplot(x="age",y="charges",data=df500,order=1,ci=None)
```

```
Out[14]: <seaborn.axisgrid.FacetGrid at 0x12592c25f90>
```



```
In [15]: df500.fillna(method='ffill',inplace=True)
x=np.array(df500['age']).reshape(-1,1)
y=np.array(df500['charges']).reshape(-1,1)
df500.dropna(inplace=True)
x_train,x_test,y_train,y_test=train_test_split(x,y,test_size=0.25)
regr=LinearRegression()
regr.fit(x_train,y_train)
print("Regression:",regr.score(x_test,y_test))
y_pred=regr.predict(x_test)
plt.scatter(x_test,y_test,color='r')
plt.plot(x_test,y_pred,color='g')
plt.show()
```

Regression: 0.09329384184108624




```
In [16]: from sklearn.linear_model import LinearRegression
from sklearn.metrics import r2_score
model=LinearRegression()
model.fit(x_train,y_train)
y_pred=model.predict(x_test)
r2=r2_score(y_test,y_pred)
print("R2 Score:",r2)
```

R2 Score: 0.09329384184108624

```
In [17]: df.isnull().sum()
```

```
Out[17]: age      0
sex        0
bmi        0
children   0
smoker     0
region     0
charges    0
dtype: int64
```

Implementon of Ridgeand Lasso Regression model

```
In [22]: from sklearn.linear_model import Lasso,Ridge
from sklearn.preprocessing import StandardScaler
```

```
In [23]: convert={"sex":{"male":1,"female":2}}
df=df.replace(convert)
df
```

```
Out[23]:
```

	age	sex	bmi	children	smoker	region	charges
0	19	2	27.900	0	yes	southwest	16884.92400
1	18	1	33.770	1	no	southeast	1725.55230
2	28	1	33.000	3	no	southeast	4449.46200
3	33	1	22.705	0	no	northwest	21984.47061
4	32	1	28.880	0	no	northwest	3866.85520
...
1333	50	1	30.970	3	no	northwest	10600.54830
1334	18	2	31.920	0	no	northeast	2205.98080
1335	18	2	36.850	0	no	southeast	1629.83350
1336	21	2	25.800	0	no	southwest	2007.94500
1337	61	2	29.070	0	yes	northwest	29141.36030

1338 rows × 7 columns

```
In [24]: convert={"smoker":{"yes":1,"no":2}}
df=df.replace(convert)
df
```

```
Out[24]:
```

	age	sex	bmi	children	smoker	region	charges
0	19	2	27.900	0	1	southwest	16884.92400
1	18	1	33.770	1	2	southeast	1725.55230
2	28	1	33.000	3	2	southeast	4449.46200
3	33	1	22.705	0	2	northwest	21984.47061
4	32	1	28.880	0	2	northwest	3866.85520
...
1333	50	1	30.970	3	2	northwest	10600.54830
1334	18	2	31.920	0	2	northeast	2205.98080
1335	18	2	36.850	0	2	southeast	1629.83350
1336	21	2	25.800	0	2	southwest	2007.94500
1337	61	2	29.070	0	1	northwest	29141.36030

1338 rows × 7 columns

```
In [25]: convert={"region":{"southeast":3,"southwest":4,"northeast":5,"northwest":6}}
df=df.replace(convert)
df
```

```
Out[25]:
```

	age	sex	bmi	children	smoker	region	charges
0	19	2	27.900	0	1	4	16884.92400
1	18	1	33.770	1	2	3	1725.55230
2	28	1	33.000	3	2	3	4449.46200
3	33	1	22.705	0	2	6	21984.47061
4	32	1	28.880	0	2	6	3866.85520
...
1333	50	1	30.970	3	2	6	10600.54830
1334	18	2	31.920	0	2	5	2205.98080
1335	18	2	36.850	0	2	3	1629.83350
1336	21	2	25.800	0	2	4	2007.94500
1337	61	2	29.070	0	1	6	29141.36030

1338 rows × 7 columns

```
In [27]: features = df.columns[0:1]
target = df.columns[-1]
#X and y values
X = df[features].values
y = df[target].values
#split
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.3, random_state=17)
print("The dimension of x_train is {}".format(X_train.shape))
print("The dimension of x_test is {}".format(X_test.shape))
#Scale features
scaler = StandardScaler()
X_train = scaler.fit_transform(X_train)
X_test = scaler.transform(X_test)
```

The dimension of x_train is (936, 1)

The dimension of x_test is (402, 1)

```
In [28]: ridgeReg=Ridge(alpha=10)
ridgeReg.fit(X_train,y_train)
train_score_ridge=ridgeReg.score(X_train,y_train)
test_score_ridge=ridgeReg.score(X_test,y_test)
print("\nRidge Model:\n")
print("The train score for ridge model is {}".format(train_score_ridge))
print("The test score for ridge model is {}".format(test_score_ridge))
```

Ridge Model:

The train score for ridge model is 0.07446228994221393

The test score for ridge model is 0.10855133360950642

```
In [29]: lr = LinearRegression()
#Fit model
lr.fit(X_train, y_train)
#predict
#prediction = lr.predict(X_test)
#actual
actual = y_test
train_score_lr = lr.score(X_train, y_train)
test_score_lr = lr.score(X_test, y_test)
print("\nLinear Regression Model:\n")
print("The train score for lr model is {}".format(train_score_lr))
print("The test score for lr model is {}".format(test_score_lr))
```

Linear Regression Model:

The train score for lr model is 0.07447061146193878

The test score for lr model is 0.10891203216512224


```
In [31]: print("\nLasso Model: \n")
lasso = Lasso(alpha = 10)
lasso.fit(X_train,y_train)
train_score_ls =lasso.score(X_train,y_train)
test_score_ls =lasso.score(X_test,y_test)
print("The train score for ls model is {}".format(train_score_ls))
print("The test score for ls model is {}".format(test_score_ls))
```

Lasso Model:

The train score for ls model is 0.07446997086306062

The test score for ls model is 0.10881427793326703

```
In [33]: pd.Series(lasso.coef_, features).sort_values(ascending = True).plot(kind = "bar")  
plt.show()
```



```
In [38]: from sklearn.linear_model import LassoCV  
lasso_cv=LassoCV(alphas=[0.0001,0.001,0.01,0.1,1,10],random_state=0).fit(X_train,y_train)  
print(lasso_cv.score(X_train,y_train))  
print(lasso_cv.score(X_test,y_test))
```

```
0.07446997086306062  
0.10881427793326703
```

```
In [40]: plt.figure(figsize=(10,10))
#add plot for ridge regression
plt.plot(features,ridgeReg.coef_,alpha=0.7,linestyle='none',marker='*',markersize=5,color='red',label='Ridge;$\alpha=1')
#add plot for Lasso regression
plt.plot(lasso_cv.coef_,alpha=0.6,linestyle='none',marker='d',markersize=6,color='blue',label='Ridge;$\alpha=grid$')
#add plot for Linear model
plt.plot(features,lr.coef_,alpha=0.4,linestyle='none',marker='o',markersize=7,color='green',label='Linear Regression')
#rotate axis
plt.xticks(rotation=90)
plt.legend()
plt.title("comparison plot of Ridge,Lasso and Linear regression model")
plt.show()
```

AttributeError

Traceback (most recent call last)

Cell In[40], line 3

```
1 plt.figure(figsize=(10,10))
2 #add plot for ridge regression
----> 3 plt.plot(features,ridgeReg.coef_,alpha=0.7,linestyle='none',marker='*',markersize=5,color='red',label='Ridge;$\alpha=10$',border=7)
4 #add plot for Lasso regression
5 plt.plot(lasso_cv.coef_,alpha=0.6,linestyle='none',marker='d',markersize=6,color='blue',label='Ridge;$\alpha=grid$')
```

File ~\AppData\Local\Programs\Python\Python311\Lib\site-packages\matplotlib\pyplot.py:2812, in plot(scalex, scaley, data, *args, **kwargs)

```
2810 @_copy_docstring_and_deprecators(Axes.plot)
2811 def plot(*args, scalex=True, scaley=True, data=None, **kwargs):
-> 2812     return gca().plot(
2813         *args, scalex=scalex, scaley=scaley,
2814         **({"data": data} if data is not None else {}), **kwargs)
```



```
In [42]: from sklearn.linear_model import RidgeCV
#Ridge Cross validation
ridge_cv = RidgeCV(alphas = [0.0001, 0.001,0.01, 0.1, 1, 10]).fit(X_train, y_train)
#score
print(ridge_cv.score(X_train, y_train))
print(ridge_cv.score(X_test, y_test))
```

```
0.07446228994221393
0.10855133360950775
```

Elastic net regression

```
In [43]: from sklearn.linear_model import ElasticNet
regr=ElasticNet()
regr.fit(X,y)
print(regr.coef_)
print(regr.intercept_)
regr.score(X,y)
```

```
[257.0684655]
3191.532406056682
```

```
Out[43]: 0.08940532368214038
```

```
In [44]: y_pred_elastic=regr.predict(X_train)
```

```
In [45]: mean_squared_error=np.mean((y_pred_elastic-y_train)**2)
print("Mean Squared Error on test set",mean_squared_error)
```

```
Mean Squared Error on test set 267460995.25217086
```

Logistic Regression

```
In [46]: import pandas as pd
import numpy as np
from sklearn.linear_model import LogisticRegression
from sklearn.preprocessing import StandardScaler
```

```
In [47]: df=pd.read_csv(r"C:\Users\sweet\Downloads\insuranceex.csv")
df
```

```
Out[47]:
```

	age	sex	bmi	children	smoker	region	charges
0	19	female	27.900	0	yes	southwest	16884.92400
1	18	male	33.770	1	no	southeast	1725.55230
2	28	male	33.000	3	no	southeast	4449.46200
3	33	male	22.705	0	no	northwest	21984.47061
4	32	male	28.880	0	no	northwest	3866.85520
...
1333	50	male	30.970	3	no	northwest	10600.54830
1334	18	female	31.920	0	no	northeast	2205.98080
1335	18	female	36.850	0	no	southeast	1629.83350
1336	21	female	25.800	0	no	southwest	2007.94500
1337	61	female	29.070	0	yes	northwest	29141.36030

1338 rows × 7 columns

```
In [48]: pd.set_option('display.max_rows',1000000000)
pd.set_option('display.max_columns',1000000000)
pd.set_option('display.width',95)
```

```
In [49]: print('This DataFrame has %d Rows and %d columns'%(df.shape))
```

This DataFrame has 1338 Rows and 7 columns

```
In [50]: convert={"smoker":{"yes":1,"no":2}}
df=df.replace(convert)
df
```

```
Out[50]:
```

	age	sex	bmi	children	smoker	region	charges
0	19	female	27.900	0	1	southwest	16884.924000
1	18	male	33.770	1	2	southeast	1725.552300
2	28	male	33.000	3	2	southeast	4449.462000
3	33	male	22.705	0	2	northwest	21984.470610
4	32	male	28.880	0	2	northwest	3866.855200
5	31	female	25.740	0	2	southeast	3756.621600
6	46	female	33.440	1	2	southeast	8240.589600
7	37	female	27.740	3	2	northwest	7281.505600
8	37	male	29.830	2	2	northeast	6406.410700
9	60	female	25.840	0	2	northwest	28923.136920
10	25	male	26.220	0	2	northeast	2721.320800
11	62	female	26.290	0	1	southeast	27808.725100

```
In [51]: convert={"sex":{"male":8,"female":9}}  
df=df.replace(convert)  
df
```

```
Out[51]:
```

	age	sex	bmi	children	smoker	region	charges
0	19	9	27.900	0	1	southwest	16884.924000
1	18	8	33.770	1	2	southeast	1725.552300
2	28	8	33.000	3	2	southeast	4449.462000
3	33	8	22.705	0	2	northwest	21984.470610
4	32	8	28.880	0	2	northwest	3866.855200
5	31	9	25.740	0	2	southeast	3756.621600
6	46	9	33.440	1	2	southeast	8240.589600
7	37	9	27.740	3	2	northwest	7281.505600
8	37	8	29.830	2	2	northeast	6406.410700
9	60	9	25.840	0	2	northwest	28923.136920
10	25	8	26.220	0	2	northeast	2721.320800
11	62	9	26.290	0	1	southeast	27808.725100

```
In [52]: features_matrix=df.iloc[:,0:4]
```

```
In [53]: target_vector=df.iloc[:, -3]
```

```
In [55]: print('The Features Matrix Has %d Rows And %d Column(s)'%(features_matrix.shape))
```

The Features Matrix Has 1338 Rows And 4 Column(s)

```
In [57]: print('The Target Matrix Has %d Rows And %d Column(s)'%(np.array(target_vector).reshape(-1,1).shape))
```

The Target Matrix Has 1338 Rows And 1 Column(s)

```
In [58]: features_matrix_standardized=StandardScaler().fit_transform(features_matrix)
```

```
In [60]: algorithm=LogisticRegression(penalty='l2',dual=False,tol=1e-4,C=1.0,fit_intercept=True,intercept_scaling=1,class_weight
```

```
In [62]: Logistic_Regression_Model=algorithm.fit(features_matrix_standardized,target_vector)
```

```
In [63]: observation=[[1,0,0.99539,-0.05889,]]
```

```
In [65]: predictions=Logistic_Regression_Model.predict(observation)
print('The Model Predicted The Observation To Belong To Class %s'%(predictions))
```

The Model Predicted The Observation To Belong To Class [2]

```
In [67]: print('The algorithm was trained to predict one of the two classes: %s'%(algorithm.classes_))
```

The algorithm was trained to predict one of the two classes: [1 2]

```
In [85]: print(" " "The model says the probability of the observation we passed belonging to class[0] Is %s" " "%(algorithm.pre
```

The model says the probability of the observation we passed belonging to class[0] Is 0.1942921563693959

```
In [86]: print(" " "The model says the probability of the observation we passed belonging to class['1'] Is %s" " "%(algorithm.p
```

The model says the probability of the observation we passed belonging to class['1'] Is 0.1942921563693959

```
In [87]: x=np.array(df['age']).reshape(-1,1)
y=np.array(df['smoker']).reshape(-1,1)
```

```
In [88]: lerg=LogisticRegression()  
         lerg.fit(x,y)  
         print(lerg.score(x,y))
```

0.7952167414050823

C:\Users\sweet\AppData\Local\Programs\Python\Python311\Lib\site-packages\sklearn\utils\validation.py:1143: DataConversionWarning: A column-vector y was passed when a 1d array was expected. Please change the shape of y to (n_samples,), for example using ravel().
y = column_or_1d(y, warn=True)

decision tree regression

```
In [89]: import numpy as np  
         import pandas as pd  
         import seaborn as sns  
         from sklearn.model_selection import train_test_split  
         from sklearn.tree import DecisionTreeClassifier
```

```
In [90]: df=pd.read_csv(r"C:\Users\sweet\Downloads\insuranceex.csv")
df
```

```
Out[90]:
```

	age	sex	bmi	children	smoker	region	charges
0	19	female	27.900	0	yes	southwest	16884.924000
1	18	male	33.770	1	no	southeast	1725.552300
2	28	male	33.000	3	no	southeast	4449.462000
3	33	male	22.705	0	no	northwest	21984.470610
4	32	male	28.880	0	no	northwest	3866.855200
5	31	female	25.740	0	no	southeast	3756.621600
6	46	female	33.440	1	no	southeast	8240.589600
7	37	female	27.740	3	no	northwest	7281.505600
8	37	male	29.830	2	no	northeast	6406.410700
9	60	female	25.840	0	no	northwest	28923.136920
10	25	male	26.220	0	no	northeast	2721.320800
11	62	female	26.290	0	ves	southeast	27808.725100

```
In [91]: df['region'].value_counts()
```

```
Out[91]: region
southeast    364
southwest    325
northwest    325
northeast    324
Name: count, dtype: int64
```

```
In [92]: df['bmi'].value_counts()
```

```
Out[92]: bmi
32.300    13
28.310     9
30.495     8
30.875     8
31.350     8
30.800     8
34.100     8
28.880     8
33.330     7
35.200     7
25.800     7
32.775     7
27.645     7
32.110     7
38.060     7
25.460     7
30.590     7
27.360     7
24.330     7
```



```
In [93]: convert={"sex":{"male":1,"female":0}}  
df=df.replace(convert)  
df
```

175	63	0	37.700	0	yes	southwest	48824.450000
176	38	1	27.835	2	no	northwest	6455.862650
177	54	1	29.200	1	no	southwest	10436.096000
178	46	0	28.900	2	no	southwest	8823.279000
179	41	0	33.155	3	no	northeast	8538.288450
180	58	1	28.595	0	no	northwest	11735.879050
181	18	0	38.280	0	no	southeast	1631.821200
182	22	1	19.950	3	no	northeast	4005.422500
183	44	0	26.410	0	no	northwest	7419.477900
184	44	1	30.690	2	no	southeast	7731.427100
185	36	1	41.895	3	yes	northeast	43753.337050
186	26	0	29.920	2	no	southeast	3981.976800
187	30	0	30.900	3	no	southwest	5325.651000

```
In [94]: x=["bmi","children"]  
y=["yes","no"]  
all_inputs=df[x]  
all_classes=df["sex"]
```

```
In [96]: (x_train,x_test,y_train,y_test)=train_test_split(all_inputs,all_classes,test_size=0.03)
```

```
In [97]: clf=DecisionTreeClassifier(random_state=0)
```

```
In [98]: clf.fit(x_train,y_train)
```

```
Out[98]: DecisionTreeClassifier(random_state=0)
```

**In a Jupyter environment, please rerun this cell to show the HTML representation or trust the notebook.
On GitHub, the HTML representation is unable to render, please try loading this page with nbviewer.org.**

```
In [99]: score=clf.score(x_test,y_test)
print(score)
```

```
0.4878048780487805
```

random forest

```
In [100]: import pandas as pd
import numpy as ny
import matplotlib.pyplot as plt,seaborn as sns
```

```
In [101]: df=pd.read_csv(r"C:\Users\sweet\Downloads\insuranceex.csv")
df
```

```
Out[101]:
```

	age	sex	bmi	children	smoker	region	charges
0	19	female	27.900	0	yes	southwest	16884.924000
1	18	male	33.770	1	no	southeast	1725.552300
2	28	male	33.000	3	no	southeast	4449.462000
3	33	male	22.705	0	no	northwest	21984.470610
4	32	male	28.880	0	no	northwest	3866.855200
5	31	female	25.740	0	no	southeast	3756.621600
6	46	female	33.440	1	no	southeast	8240.589600
7	37	female	27.740	3	no	northwest	7281.505600
8	37	male	29.830	2	no	northeast	6406.410700
9	60	female	25.840	0	no	northwest	28923.136920
10	25	male	26.220	0	no	northeast	2721.320800
11	62	female	26.290	0	ves	southeast	27808.725100

```
In [102]: df['charges'].value_counts()
```

```
Out[102]: charges
1639.563100    2
16884.924000   1
29330.983150   1
2221.564450    1
19798.054550   1
13063.883000   1
13555.004900   1
44202.653600   1
10422.916650   1
7243.813600    1
11945.132700   1
6311.952000    1
1682.597000    1
5272.175800    1
27218.437250   1
19719.694700   1
4877.981050    1
46255.112500   1
3535.050650    1
```

```
In [103]: m={"region":{"southeast":1,"southwest":2,"northeast":3,"northwest":4}}
df=df.replace(m)
print(df)
```

122	20	female	28.975	0	no	4	2257.475250
123	44	male	31.350	1	yes	3	39556.494500
124	47	female	33.915	3	no	4	10115.008850
125	26	female	28.785	0	no	3	3385.399150
126	19	female	28.300	0	yes	2	17081.080000
127	52	female	37.400	0	no	2	9634.538000
128	32	female	17.765	2	yes	4	32734.186300
129	38	male	34.700	2	no	2	6082.405000
130	59	female	26.505	0	no	3	12815.444950
131	61	female	22.040	0	no	3	13616.358600
132	53	female	35.900	2	no	2	11163.568000
133	19	male	25.555	0	no	4	1632.564450
134	20	female	28.785	0	no	3	2457.211150
135	22	female	28.050	0	no	1	2155.681500
136	19	male	34.100	0	no	2	1261.442000
137	22	male	25.175	0	no	4	2045.685250
138	54	female	31.900	3	no	1	27322.733860
139	22	female	36.000	0	no	2	2166.732000
140	34	male	22.420	2	no	3	27375.904780
141	26	male	32.490	1	no	3	3490.549100

```
In [106]: df.shape
```

```
Out[106]: (1338, 7)
```

```
In [107]: from sklearn.ensemble import RandomForestClassifier
rfc=RandomForestClassifier()
rfc.fit(x_train,y_train)
```

```
Out[107]: RandomForestClassifier()
```

**In a Jupyter environment, please rerun this cell to show the HTML representation or trust the notebook.
On GitHub, the HTML representation is unable to render, please try loading this page with nbviewer.org.**

```
In [108]: rf=RandomForestClassifier()
```

```
In [110]: params={'max_depth':[2,3,5,10,20], 'min_samples_leaf':[5,10,20,50,100,200], 'n_estimators':[10,25,30,50,100,200]}
```

```
In [113]: from sklearn.model_selection import GridSearchCV
grid_search=GridSearchCV(estimator=rf,param_grid=params,cv=2,scoring="accuracy")
grid_search.fit(x_train,y_train)
```

```
Out[113]: GridSearchCV(cv=2, estimator=RandomForestClassifier(),
                      param_grid={'max_depth': [2, 3, 5, 10, 20],
                                   'min_samples_leaf': [5, 10, 20, 50, 100, 200],
                                   'n_estimators': [10, 25, 30, 50, 100, 200]},
                      scoring='accuracy')
```

**In a Jupyter environment, please rerun this cell to show the HTML representation or trust the notebook.
On GitHub, the HTML representation is unable to render, please try loading this page with nbviewer.org.**

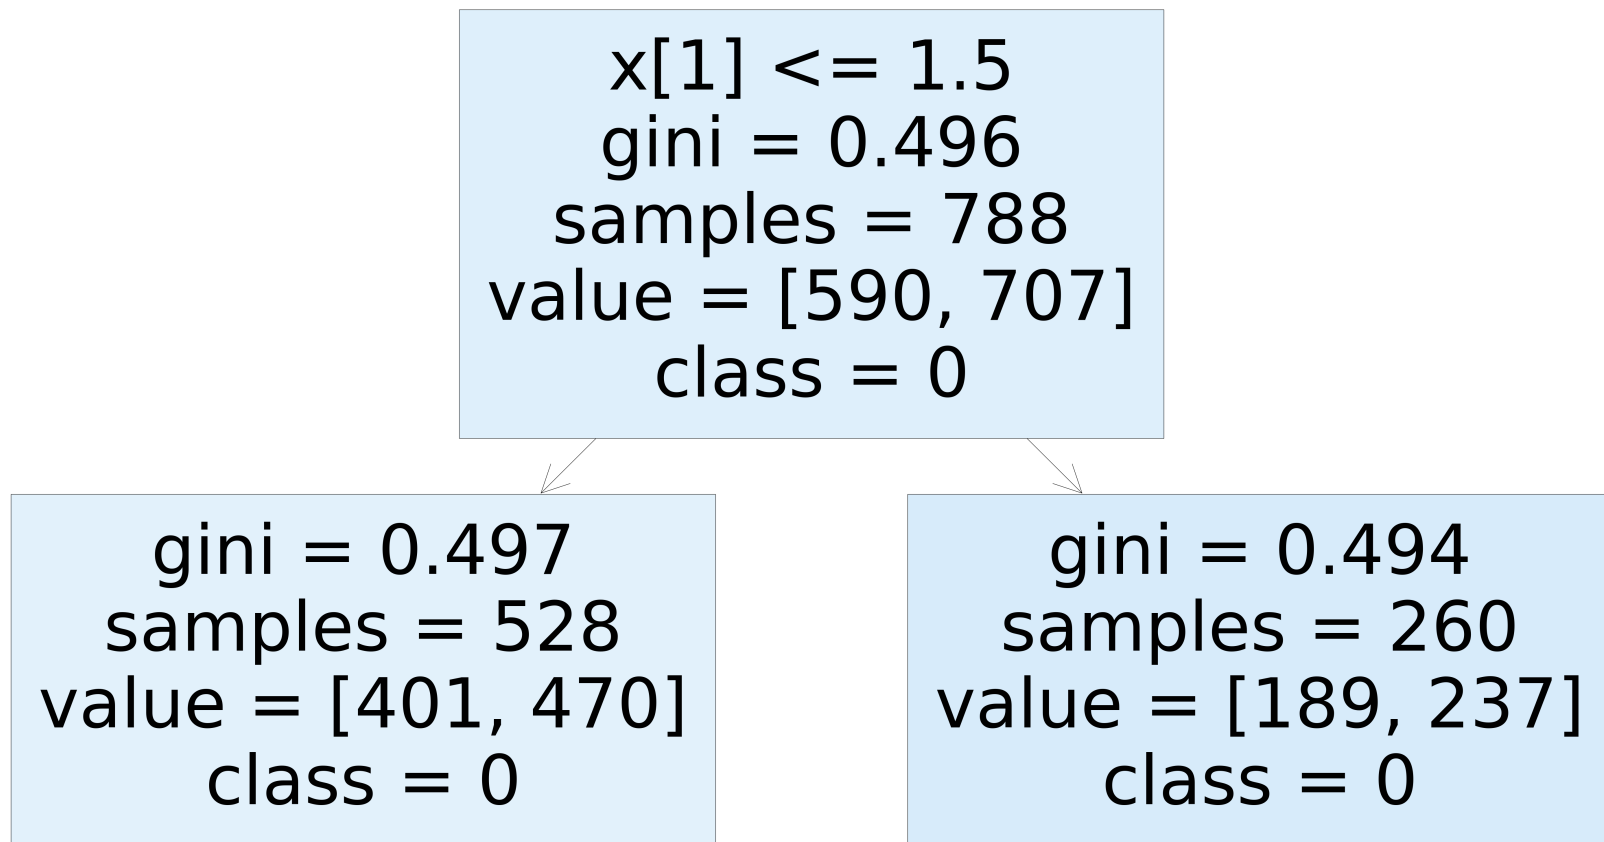
```
In [115]: grid_search.best_score_
```

```
Out[115]: 0.5219628012707109
```

```
In [116]: rf_best=grid_search.best_estimator_
print(rf_best)
```

```
RandomForestClassifier(max_depth=20, min_samples_leaf=200, n_estimators=10)
```

```
In [117]: from sklearn.tree import plot_tree  
plt.figure(figsize=(80,40))  
plot_tree(rf_best.estimators_[4],class_names=['1','0'],filled=True);
```



```
In [119]: rf_best.feature_importances_
```

```
Out[119]: array([0.47848867, 0.52151133])
```

```
In [122]: score=rfc.score(x_test,y_test)  
          print(score)
```

```
0.4146341463414634
```